

# Machine Learning

## Kernel Methods

Alberto Maria Metelli and Francesco Trovò

# Definition of Different Models

What to do in the case the model you are considering is not performing well even by tuning properly the parameters (cross-validation)?

We have two opposite options:

- simplify the model
- **increase its complexity**

In the second hypothesis, one might see the problem in a more complex space:

- use handcrafted features
- look at the problem in the **kernel** space

# Constructing Kernels

# Constructing Kernels

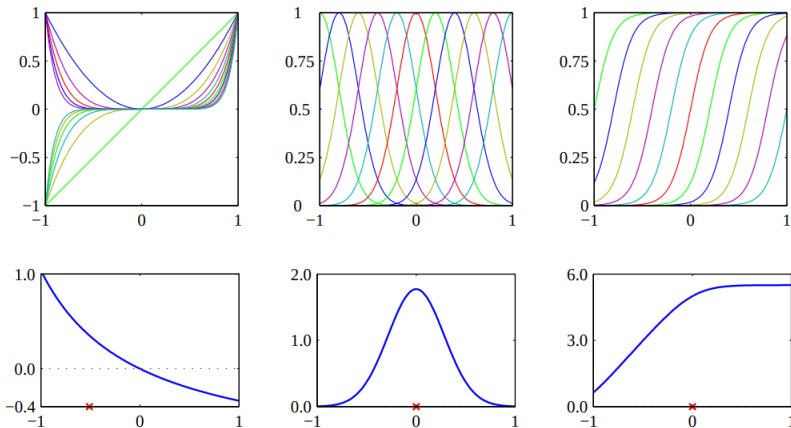
- To exploit kernel substitution need **valid** kernel functions
- First method
  - **Choose a feature space** mapping  $\phi(\mathbf{x})$  and use it to find corresponding kernel

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}') = \sum_{i=1}^M \phi_i(\mathbf{x}) \phi_i(\mathbf{x}')$$

- where  $\phi(\mathbf{x})$  are basis functions such as polynomial
- for each  $i$  we choose  $\phi_i(x) = x^i$  in the one-dimensional case

# Construction of Kernel Functions from Basis Functions

One-dimensional input space



# Testing whether a function is a valid kernel

- Second method
  - Without having to construct the function  $\phi(\mathbf{x})$  explicitly
- **Necessary and sufficient** condition for a function  $k(\mathbf{x}, \mathbf{x}')$  to be a kernel is
  - Gram matrix  $\mathbf{K}$ , whose elements are given by  $k(\mathbf{x}_n, \mathbf{x}_m)$  is positive semi-definite for all possible choices of the set  $\{\mathbf{x}_n\}$ 
    - Positive semi-definite is **not** the same thing as a matrix whose elements are non-negative
    - It means  $\mathbf{y}^\top \mathbf{K} \mathbf{y} \geq 0$  for non-zero vectors  $\mathbf{y}$  with real entries, i.e.,  $\sum_n \sum_m K_{n,m} y_n y_m \geq 0$  for any real numbers  $y$

## Theorem

*Mercer's theorem **Any** continuous, symmetric, positive semi-definite kernel function  $k(x, x')$  can be expressed as a **dot product** in a high-dimensional space*

- **New kernels** can be constructed from simpler kernels as **building blocks**

# Techniques for Constructing Kernels

Given valid kernels  $k_1(\mathbf{x}, \mathbf{x}')$  and  $k_2(\mathbf{x}, \mathbf{x}')$  the following new kernels will be valid

- ❶  $k(\mathbf{x}, \mathbf{x}') = ck_1(\mathbf{x}, \mathbf{x}')$
- ❷  $k(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})k_1(\mathbf{x}, \mathbf{x}')f(\mathbf{x}')$ , where  $f(\cdot)$  is any function
- ❸  $k(\mathbf{x}, \mathbf{x}') = q(k_1(\mathbf{x}, \mathbf{x}'))$ , where  $q(\cdot)$  is a polynomial with non-negative coefficients
- ❹  $k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}'))$
- ❺  $k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$
- ❻  $k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}')$
- ❼  $k(\mathbf{x}, \mathbf{x}') = k_3(\phi(\mathbf{x}), \phi(\mathbf{x}'))$ , where  $\phi(\mathbf{x})$  is a function from  $\mathbf{x}$  to  $\mathbb{R}^M$
- ❽  $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T A \mathbf{x}'$ , where  $A$  is a symmetric positive semidefinite matrix
- ❾  $k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a) + k_b(\mathbf{x}_b, \mathbf{x}'_b)$ , where  $x_a$  and  $x_b$  are variables with  $\mathbf{x} = (x_a, x_b)$
- ❿  $k(\mathbf{x}, \mathbf{x}') = k_a(\mathbf{x}_a, \mathbf{x}'_a)k_b(\mathbf{x}_b, \mathbf{x}'_b)$

# Gaussian Processes



# GP Definition

- A Gaussian process is defined as a probability distribution over functions  $y(\mathbf{x}_i)$ 
  - values of  $y(\mathbf{x}_i)$  evaluated at any set of points  $\mathbf{x}_1, \dots, \mathbf{x}_N$  jointly have a Gaussian distribution
- This distribution is completely specified by the mean and the covariance:
  - usually, we do not have any prior information about the mean of  $y(\mathbf{x})$ , so we take it to be zero
  - the covariance is given by the **kernel** function

$$\text{Cov}[y(\mathbf{x}_i), y(\mathbf{x}_j) | \mathbf{x}_i, \mathbf{x}_j] = \mathbb{E}[y(\mathbf{x}_i) y(\mathbf{x}_j) | \mathbf{x}_i, \mathbf{x}_j] = K(\mathbf{x}_i, \mathbf{x}_j)$$

- With this formulation, **Gaussian Process** (GP) are kernel methods that can be applied to solve regression problems

# Output Modeling

- The **target** is  $t = y(\mathbf{x}) + \varepsilon$ , where  $\varepsilon \sim \mathcal{N}(0, \sigma^2)$  is a noise independent on the point  $\mathbf{x}$
- The conditional distribution of the targets  $\mathbf{t}_N = (t_1, \dots, t_N)^\top$  of size  $N$  is:

$$p(\mathbf{t}_N | \mathbf{y}_N) = \mathcal{N}(\mathbf{t}_N | \mathbf{y}_N, \sigma^2 \mathbf{I}_N) \quad \text{where } \mathbf{y}_N = (y(\mathbf{x}_1), \dots, y(\mathbf{x}_N))^\top$$

- The **prior** is  $p(\mathbf{y}_N) = \mathcal{N}(\mathbf{0}, \mathbf{K}_N)$ , where:

$$\mathbf{K}_N = \begin{pmatrix} K(\mathbf{x}_1, \mathbf{x}_1) & \dots & K(\mathbf{x}_1, \mathbf{x}_N) \\ \vdots & \ddots & \vdots \\ K(\mathbf{x}_N, \mathbf{x}_1) & \dots & K(\mathbf{x}_N, \mathbf{x}_N) \end{pmatrix}$$

- Thus, the marginal distribution of the target is:

$$p(\mathbf{t}_N) = \int p(\mathbf{t}_N | \mathbf{y}_N) p(\mathbf{y}) \, d\mathbf{y}_N = \mathcal{N}(\mathbf{t}_N | \mathbf{0}, \mathbf{C}_N) \quad \text{where } \mathbf{C}_N = \mathbf{K}_N + \sigma^2 \mathbf{I}_N$$

# Making Predictions

- We want to **predict** the target  $t_{N+1}$  corresponding to a specific unseen input  $\mathbf{x}_{N+1}$
- From the definition we have:

$$p(\mathbf{t}_{N+1}) = \mathcal{N}(\mathbf{t}_{N+1} | \mathbf{0}, \mathbf{C}_{N+1}),$$

where:

$$\mathbf{C}_{N+1} = \begin{pmatrix} \mathbf{C}_N & \mathbf{k} \\ \mathbf{k}^\top & c \end{pmatrix}$$

$$\mathbf{k} = (K(\mathbf{x}_1, \mathbf{x}_{N+1}), \dots, K(\mathbf{x}_N, \mathbf{x}_{N+1}))^\top \quad c = K(\mathbf{x}_{N+1}, \mathbf{x}_{N+1}) + \sigma^2$$

- We need to compute  $p(t_{N+1} | \mathbf{t}_N, \mathbf{x}_1, \dots, \mathbf{x}_N) = \mathcal{N}(m(\mathbf{x}_{N+1}), \sigma^2(\mathbf{x}_{N+1}))$  where
  - Mean:  $m(\mathbf{x}_{N+1}) = \mathbf{k}^\top \mathbf{C}_N^{-1} \mathbf{t}$
  - Variance:  $\sigma^2(\mathbf{x}_{N+1}) = c - \mathbf{k}^\top \mathbf{C}_N^{-1} \mathbf{k}$

# GPs in Python

Model the relationship between petal length and width as a GP:

- Load the data and normalize them
- Select the values of:
  - noise variance  $\sigma^2 = \mathbb{V}\text{ar}[\varepsilon] = 0.2$
  - constant  $k = 1$
  - lengthscale  $l = 0.8$

$$K(\mathbf{x}_i, \mathbf{x}_j) = k \exp \left\{ -\frac{\|\mathbf{x}_i - \mathbf{x}_j\|_2^2}{2l^2} \right\}$$

- Initialize a GP regression model (`GaussianProcessRegressor`)
- Predict new values


# Hyperparameters

While GPs are a **non-parametric** methods,  $\sigma^2$  and the parameters of the kernel (e.g.,  $l$  and  $k$ ) has to be estimated or set:

- using **a priori** information on the problem we are analyzing
- maximizing their **log-likelihood** on an independent dataset
- possibly improved as new data are collected

Caveat: most of the time you will see that they are estimated using the same data used for the prediction. This is clearly not a good ML practice (equivalent to overfitting)



# Results on the Iris Dataset



Parameters:  $k = 3$ ,  $l = 0.8$ , and  $\sigma^2 = 0.2$

# Modify the Lengthscale



- Left:  $k = 3$ ,  $l = 8$ , and  $\sigma^2 = 0.2$
- Right:  $k = 3$ ,  $l = 0.08$ , and  $\sigma^2 = 0.2$



## Controls the smoothness of the GP

# Modify the Noise

- Left:  $k = 3$ ,  $l = 0.8$ , and  $\sigma^2 = 10$
- Right:  $k = 3$ ,  $l = 0.8$ , and  $\sigma^2 = 0.002$



Controls the target noise of the GP



# Support Vector Machines

# Support Vector Machines (SVM)

- Flexible and **theoretically supported** method
- Initially applied to classification only, over the years it has been extended to deal with regression, clustering and anomaly detection problems
- Idea: find the **hyperplane maximizing the margins** (distance between the boundary and the points)

# Support Vector Machines (SVM)

- Hypothesis space:  $y(\mathbf{x}) = f(\mathbf{x}; \mathbf{w}) = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$
- Loss function computed over  $\mathcal{D} = \{(\mathbf{x}_n, t_n)\}_{n=1}^N$  with  $t_n \in \{-1, 1\}$ :

$$\begin{array}{ll} \underset{\mathbf{w}, \zeta_1, \dots, \zeta_N}{\text{minimize}} & \frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_{n=1}^N \zeta_i \\ \text{s.t.} & t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1 - \zeta_i \quad \forall n \in \{1, \dots, N\} \\ & \zeta_i \geq 0 \quad \forall n \in \{1, \dots, N\} \end{array}$$

where  $C > 0$  is a hyperparameter

- Optimization method: sequential quadratic optimization

# Linear SVM in Python


To train a linear classification SVM:

- Define an SVM: `SVM_model.svm.SVC(kernel='linear')`
- Train the SVM: `SVM_model.fit(input, target)`

We are interested to determine:

- Boundary  $\mathbf{w}^T \mathbf{x}_n + b = 0$
- Margins  $\mathbf{w}^T \mathbf{x}_n + b = \pm 1$
- Support vectors (`SVM_model.support_vectors_`)

# Results on the Iris Dataset



# Adding a Kernel

The use of kernels in the SVM is almost native (non-parametric method):

- Hypothesis space:  $y(\mathbf{x}) = f(\mathbf{x}; \mathbf{w}) = \text{sign} \left( \sum_{n=1}^N \alpha_n t_n K(\mathbf{x}_n, \mathbf{x}) + b \right)$
- Loss measure: loss function in the dual formulation
- Optimization method: quadratic optimization

In Python:

- Define an SVM: `SVM_model.svm.SVC()`
- Train the SVM: `SVM_model.fit(input, target)`

We do not have an explicit formula for the boundary and the margins anymore

# Results on the Iris Dataset

